

final

May 15, 2023

# 1 Analysing dialogues between speakers of different ages

## 1.1 Task 1: Description of Data and Methods (10/50) – (~max 600 words)

BNC2014spoken-xml v1.1 dataset contains the following files that are used for this project:

- spoken
  - tagged : all the tagged data in xml format. Contains individual words for each utterances and their attributes.
  - untgged : all the untagged data in xml format. Contains utterances and speakers for that particualr dialogues.
  - metadata : text and XML files containing corpus metadata

Q1. What basic preprocessing steps would be needed to work with this data (e.g. tokenization, embeddings etc)

Ans: The basic preprocessing steps that would be needed to work with this data are:

- Data Gathering : The dataset is in xml format and is split into tagged and untagged data. We need to gather the data that is required for our analysis and store it in a dataframe.
- Data Cleaning : The data contains a lot of unnecessary information that is not required for our analysis. We need to clean the data and remove the unnecessary information.
- Data Preprocessing : The data contains a lot of noise and we need to remove the noise from the data. We need to remove the stop words, punctuations, special characters, etc. from the data. This contains the tokenization of the data and extracting the features from the data.

Q2. What feature extraction or analysis methods can you use to gain insights into the data before modeling? (basic counts like sentence length, proposition of pos tags, using linguistic resources, or clustering)

Ans: What I did was I am collecting the utterance from the xml file and calculating a list of linguistic features for each utterance. The features that I am calculating are:

- Lexical Diversity : to check the uniqueness
- Sentence Length : to check the length of the sentence
- Discourse Markers Count : the number of discourse markers in the sentence
- Fillers Count : the number of fillers in the sentence
- Hesitation Count : the number of hesitation in the sentence
- Speech Rate : the speech rate of the sentence

All of these features are calculated in the function 'extract\_linguistic\_features' which takes a preprocessed utterance(stop words, punctuations, special characters removed) as input and returns a dictionary of linguistic features. Visualizations of these features are also included.

Q3. What are common deep learning approaches to extracting information from textual data, what challenges and trade-offs are there when making an appropriate choice?

Ans: Some of the DL approaches are Long Short Term Memory (LSTM), Recurrent Neural Network (RNN), Convolutional Neural Network (CNN), Transformer models like BERT, etc. I have also included BERT model in my analysis. The challenges are huge when it comes to DL approaches.

- Data Availability : The data that is required for DL approaches is huge and it is not always available. We need to have a huge amount of data to train the model.
- Data Quality : The data that is required for DL approaches should be of high quality. The data should be clean and should not contain any noise.
- Model Complexity : The model that is used for DL approaches is very complex and it is very difficult to understand the model.
- Computational Power : Requires a lot of computational power to train the model.

and many more.

Q4. How would you evaluate machine learning systems that generate or classify text data where the gold standard can vary? or, how can you examine errors? (e.g. metrics of fluency, faithfulness, or use of confusion matrices, precision, recall, F1 score, evaluating performance in different subsets of data)

Ans:

- Confusion Matrix : We can use the confusion matrix to evaluate the performance of the machine learning systems. This will help us to understand the performance of the machine learning systems and we can use this information to improve the performance of the machine learning systems.
- Precision : Precision is the ratio of correctly predicted positive observations to the total predicted positive observations.
- Recall : Recall is the ratio of correctly predicted positive observations to the all observations in actual class.
- F1 Score : F1 Score is the weighted average of Precision and Recall. Therefore, this score takes both false positives and false negatives into account.
- Accuracy : Accuracy is the most intuitive performance measure and it is simply a ratio of correctly predicted observation to the total observations.

## 1.2 Task 2: Investigate, model, and report on insights from friends dialogue data (40/50) – (~max

1800 words)

```
[1]: import os
import re
import pandas as pd
```

```

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Set the path to the BNC2014 dataset
bnc_path = "bnc2014spoken-xml"

```

```

[2]: import nltk
      nltk.download("punkt")
      nltk.download("stopwords")
      nltk.download("wordnet")

```

```

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\abhij\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\abhij\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\abhij\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!

```

```

[2]: True

```

```

[3]: dir_corpus = 'bnc2014spoken-xml/spoken/untagged/'
      print(dir_corpus)
      f_names = os.listdir(dir_corpus)
      f_paths = [f"{dir_corpus}{f_name}" for f_name in f_names]

      # f_paths = ['bnc2014spoken-xml/spoken/untagged/S2A5.xml']

      f_paths

```

```

bnc2014spoken-xml/spoken/untagged/

```

```

[3]: ['bnc2014spoken-xml/spoken/untagged/S23A.xml',
      'bnc2014spoken-xml/spoken/untagged/S24A.xml',
      'bnc2014spoken-xml/spoken/untagged/S24D.xml',
      'bnc2014spoken-xml/spoken/untagged/S24E.xml',
      'bnc2014spoken-xml/spoken/untagged/S263.xml',
      'bnc2014spoken-xml/spoken/untagged/S26N.xml',
      'bnc2014spoken-xml/spoken/untagged/S27D.xml',
      'bnc2014spoken-xml/spoken/untagged/S28F.xml',
      'bnc2014spoken-xml/spoken/untagged/S29Q.xml',
      'bnc2014spoken-xml/spoken/untagged/S29X.xml',

```

```
'bnc2014spoken-xml/spoken/untagged/STK7.xml',
'bnc2014spoken-xml/spoken/untagged/STKH.xml',
'bnc2014spoken-xml/spoken/untagged/STKV.xml',
...]
```

```
[4]: dialogue_data = []
for path in f_paths: # for each of the dialogues
    print(path)
    df_utts = pd.read_xml(path, xpath="//u")
    df_speakers = pd.read_xml(path, xpath="//speaker")
    dialogue_data.append((df_utts, df_speakers))
    # break
```

```
bnc2014spoken-xml/spoken/untagged/S23A.xml
bnc2014spoken-xml/spoken/untagged/S24A.xml
bnc2014spoken-xml/spoken/untagged/S24D.xml
bnc2014spoken-xml/spoken/untagged/S24E.xml
bnc2014spoken-xml/spoken/untagged/S263.xml
bnc2014spoken-xml/spoken/untagged/S26N.xml
bnc2014spoken-xml/spoken/untagged/S27D.xml
bnc2014spoken-xml/spoken/untagged/S28F.xml
bnc2014spoken-xml/spoken/untagged/S29Q.xml
bnc2014spoken-xml/spoken/untagged/S29X.xml
bnc2014spoken-xml/spoken/untagged/S2A5.xml
bnc2014spoken-xml/spoken/untagged/S2AJ.xml
bnc2014spoken-xml/spoken/untagged/S2AX.xml
bnc2014spoken-xml/spoken/untagged/S2B5.xml
bnc2014spoken-xml/spoken/untagged/S2C9.xml
bnc2014spoken-xml/spoken/untagged/S2CY.xml
bnc2014spoken-xml/spoken/untagged/S2DD.xml
bnc2014spoken-xml/spoken/untagged/S2E2.xml
bnc2014spoken-xml/spoken/untagged/S2EF.xml
bnc2014spoken-xml/spoken/untagged/S2FQ.xml
bnc2014spoken-xml/spoken/untagged/S2FT.xml
bnc2014spoken-xml/spoken/untagged/S2GC.xml
bnc2014spoken-xml/spoken/untagged/S2GS.xml
bnc2014spoken-xml/spoken/untagged/S2JK.xml
bnc2014spoken-xml/spoken/untagged/S2JV.xml
bnc2014spoken-xml/spoken/untagged/S2K6.xml
bnc2014spoken-xml/spoken/untagged/S2K7.xml
bnc2014spoken-xml/spoken/untagged/S2KP.xml
bnc2014spoken-xml/spoken/untagged/S2LC.xml
bnc2014spoken-xml/spoken/untagged/S2LD.xml
bnc2014spoken-xml/spoken/untagged/S2NQ.xml
bnc2014spoken-xml/spoken/untagged/S2PS.xml
bnc2014spoken-xml/spoken/untagged/S2PY.xml
bnc2014spoken-xml/spoken/untagged/S2QU.xml
```

```

bnc2014spoken-xml/spoken/untagged/SZLE.xml
bnc2014spoken-xml/spoken/untagged/SZME.xml
bnc2014spoken-xml/spoken/untagged/SZNA.xml
bnc2014spoken-xml/spoken/untagged/SZNG.xml
bnc2014spoken-xml/spoken/untagged/SZNP.xml
bnc2014spoken-xml/spoken/untagged/SZP6.xml
bnc2014spoken-xml/spoken/untagged/SZPS.xml
bnc2014spoken-xml/spoken/untagged/SZQ9.xml
bnc2014spoken-xml/spoken/untagged/SZQX.xml
bnc2014spoken-xml/spoken/untagged/SZR7.xml
bnc2014spoken-xml/spoken/untagged/SZRJ.xml
bnc2014spoken-xml/spoken/untagged/SZT4.xml
bnc2014spoken-xml/spoken/untagged/SZVB.xml
bnc2014spoken-xml/spoken/untagged/SZVC.xml
bnc2014spoken-xml/spoken/untagged/SZW4.xml
bnc2014spoken-xml/spoken/untagged/SZXQ.xml
bnc2014spoken-xml/spoken/untagged/SZYV.xml

```

```

[5]: # Concatenate all speakers dataframes into one
df_all_speakers = pd.concat([speaker_data[1] for speaker_data in dialogue_data])
df_all_speakers

```

```

[5]:      id exactage age1994 agerange gender          nat
0   S0021      27   25_34   19_29      F      British \
1   S0032      28   25_34   19_29      M      British
2   S0094      33   25_34   30_39      F      British
3   S0095      33   25_34   30_39      M      British
0   S0261      41   35_44   40_49      M  British/New Zealand
..   ...      ...      ...      ...      ...      ...
1   S0510      47   45_59   40_49      F      British
0   S0058      23   15_24   19_29      F      British
1   S0120      23   15_24   19_29      M  British & German
0   S0428      27   25_34   19_29      F      British
1   S0432      23   15_24   19_29      F      White British

      birthplace birthcountry    l1  lingorig ...
0      Swindon      England  English  England ... \
1      Yoevil      England  English  England ...
2      Swindon      England  English  England ...
3      Camarthen  Scotland  English  England ...
0      Wellington  New Zealand  English  England/NZ ...
..      ...      ...      ...      ...
1      England      England  English  England ...
0  Sunderland, Tyne and Wear      England  English  England ...
1      Pembury, Kent      England  English  England ...
0  Aylesbury, Buckinghamshire      England  English  England ...
1      Lincoln      England  English  England ...

```

	dialect_l2	dialect_l3	dialect_l4	edqual	
0	england	south	southwest	5_postgrad	\
1	england	south	southwest	4_graduate	
2	england	south	southwest	5_postgrad	
3	wales	wales	wales	5_postgrad	
0	non_uk	non_uk	non_uk	4_graduate	
..	...	...	...	...	
1	england	south	unspecified	5_postgrad	
0	england	north	northeast	4_graduate	
1	england	south	unspecified	5_postgrad	
0	england	south	eastern_engl	5_postgrad	
1	england	unspecified	unspecified	4_graduate	

	occupation	socgrade	nssec	l2	
0	Teacher	B	2	None	\
1	Software developer	A	1_2	None	
2	PhD student	A	1_2	German	
3	Self employed maker	E	uncat	None	
0	Entrepreneur	A	1_2	NaN	
..	...	...	...	...	
1	Receptionist	D	6	NaN	
0	Corpus Administrator	B	2	NaN	
1	Graduate Civil Engineer	C1	4	NaN	
0	Language Research Co-ordinator	A	1_2	NaN	
1	Language Research Administrator	A	1_2	NaN	

	fls	in_core
0	None	y
1	None	y
2	Welsh -- Beginner	y
3	None	y
0	None	n
..	...	...
1	NaN	y
0	None	n
1	German -- Advanced; French -- Advanced	n
0	Spanish -- level unspecified; Italian -- level...	y
1	Spanish -- level unspecified; Chinese -- level...	n

[3593 rows x 25 columns]

```
[6]: # Concatenate all utterances dataframes into one
df_all_utts = pd.concat([utt_data[0] for utt_data in dialogue_data])
df_all_utts
```

```
[6]:      n      who      u
0      1  S0094      words \
1      2  S0095  it's a games word? like a computer games word?
2      3  S0032      yeah yeah
3      4  S0095      oh
4      5  S0032      I it's something I
..    ...    ...
315  316  S0432      but
316  317  S0428      None
317  318  S0432      I'll just stick with it yeah
318  319  S0428      None
319  320  S0432      mm
```

```
      unclear      trans whoConfidence      vocal foreign      anon      pause
0      None      None      None      NaN      None      NaN      NaN \
1      None      None      None      NaN      None      NaN      NaN
2      None      None      None      NaN      None      NaN      NaN
3      oh that's nice      None      None      NaN      None      NaN      NaN
4      have really heard      overlap      None      NaN      None      NaN      NaN
..    ...    ...    ...    ...    ...    ...
315      Mai Li      None      None      NaN      NaN      NaN      NaN
316      Sha Li      None      None      NaN      NaN      NaN      NaN
317      None      overlap      None      NaN      NaN      NaN      NaN
318      None      overlap      None      NaN      NaN      NaN      NaN
319      None      overlap      None      NaN      NaN      NaN      NaN
```

```
      trunc      shift      event
0      None      NaN      NaN
1      None      NaN      NaN
2      None      NaN      NaN
3      None      NaN      NaN
4      None      NaN      NaN
..    ...    ...    ...
315  None      NaN      NaN
316  None      NaN      NaN
317  let      NaN      NaN
318  None      NaN      NaN
319  None      NaN      NaN
```

[1248110 rows x 13 columns]

```
[7]: def get_speaker_age(speakerid):
      # get the age of the speaker with the given speakerid
      age = df_all_speakers[df_all_speakers['id'] == speakerid]['agerange'].values
      if len(age) == 0:
          return None
      else:
```

```

        return age[0]

get_speaker_age('S0432')

```

```
[7]: '19_29'
```

```

[8]: def preprocess_utterance(utterance):
    if utterance is None:
        return ''
    # Remove punctuation and numbers from the conversation transcript
    utterance = re.sub(r'[^\w\s]', '', utterance) # Remove punctuation
    utterance = re.sub(r'\d+', '', utterance) # Remove numbers
    utterance = utterance.strip() # Remove leading/trailing whitespaces
    return utterance

preprocess_utterance("yeah I know but that's why we're an hour late isn't it?")

```

```
[8]: 'yeah I know but thats why were an hour late isnt it'
```

```

[9]: from nltk.tokenize import word_tokenize

def extract_linguistic_features(preprocessed_conversation):
    # Extract linguistic features from the preprocessed conversation
    if preprocessed_conversation is None:
        return ""

    # Vocabulary Features
    tokens = word_tokenize(preprocessed_conversation)

    word_freq = nltk.FreqDist(tokens)
    lexical_diversity = len(word_freq) / len(tokens) if len(tokens) > 0 else 0
    # specific_word_usage = word_freq['specific_word'] if 'specific_word' in
    ↪word_freq else 0

    # Syntactic Features
    sentences = nltk.sent_tokenize(preprocessed_conversation)
    sentence_length = sum(len(word_tokenize(sentence)) for sentence in
    ↪sentences) / len(sentences) if len(sentences) > 0 else 0

    # Discourse Features
    discourse_markers = ['like', 'you know', 'basically']
    discourse_marker_count = sum(preprocessed_conversation.lower().
    ↪count(marker) for marker in discourse_markers)
    fillers = ['uh', 'um', 'hmm', 'huh', 'er', 'ah', 'eh', 'mm']
    filler_count = sum(preprocessed_conversation.lower().count(filler) for
    ↪filler in fillers)
    hesitation_count = discourse_marker_count + filler_count

```



```

# Speech Rate
words = word_tokenize(preprocessed_conversation)
speech_rate = len(words) / len(sentences) if len(sentences) > 0 else 0

# Return the extracted linguistic features as a dictionary
linguistic_features = {
    'lexical_diversity': lexical_diversity, # each word is unique
    # 'specific_word_usage': specific_word_usage, # use of specific words
    'sentence_length': sentence_length,
    'discourse_marker_count': discourse_marker_count,
    'filler_count': filler_count,
    'hesitation_count': hesitation_count,
    'speech_rate': speech_rate
}

# Convert the dictionary to a string representation
features_string = " ".join([f"{key}:{value}" for key, value in
↪linguistic_features.items()])

return features_string, linguistic_features

extract_linguistic_features(preprocess_utterance("yeah I know but that's why
↪we're an hour late isn't it?"))[0]

```

```

[9]: 'lexical_diversity:1.0 sentence_length:12.0 discourse_marker_count:0
filler_count:2 hesitation_count:2 speech_rate:12.0'

```

```

[10]: # loop through each df_all_speakers and get the age of each speaker
# Initialize lists to store speaker information and extracted features

speaker_ids = []
age_labels = []
utterances = []
linguistic_features = []

for i in df_all_utts.index:
    speaker_id = df_all_utts["who"].iloc[i]
    age = get_speaker_age(speaker_id)

    utterance = df_all_utts["u"].iloc[i]
    preprocessed_utterance = preprocess_utterance(utterance)
    # Step 3: Feature Extraction
    linguistic_feature = extract_linguistic_features(preprocessed_utterance)[0]

    speaker_ids.append(speaker_id)
    age_labels.append(age)

```

```

utterances.append(preprocessed_utterance)
linguistic_features.append(linguistic_feature)

# Create a DataFrame to store the speaker-level data
df = pd.DataFrame({"SpeakerID": speaker_ids, "AgeLabel": age_labels,
↳ "Utterance": utterances, "Features": linguistic_features})

```

```
[11]: df
```

```

[11]:      SpeakerID AgeLabel      Utterance
0          S0094    30_39      words \
1          S0095    30_39      its a games word like a computer games word
2          S0032    19_29      yeah yeah
3          S0095    30_39      oh
4          S0032    19_29      I its something I
...
1248105      S0021    19_29      I have like
1248106      S0094    30_39
1248107      S0032    19_29
1248108      S0021    19_29  I have like tomato ketchup emergencies if I ru...
1248109      S0032    19_29  thats not as bad as the barbecue that we had t...

      Features
0  lexical_diversity:1.0 sentence_length:1.0 disc...
1  lexical_diversity:0.6666666666666666 sentence_...
2  lexical_diversity:0.5 sentence_length:2.0 disc...
3  lexical_diversity:1.0 sentence_length:1.0 disc...
4  lexical_diversity:0.75 sentence_length:4.0 dis...
...
1248105  lexical_diversity:1.0 sentence_length:3.0 disc...
1248106  lexical_diversity:0 sentence_length:0 discours...
1248107  lexical_diversity:0 sentence_length:0 discours...
1248108  lexical_diversity:0.7391304347826086 sentence_...
1248109  lexical_diversity:0.9090909090909091 sentence_...

[1248110 rows x 4 columns]

```

### 1.2.1 Visualization

```
[12]: df_vis = df.copy()
```

```


[13]: import matplotlib.pyplot as plt

# Count the occurrences of each age range category
age_counts = df_vis['AgeLabel'].value_counts()

# Create a bar plot

```


```
plt.figure(figsize=(8, 6))
plt.bar(age_counts.index, age_counts.values)
plt.xlabel('Age Range')
plt.ylabel('Count')
plt.title('Age Range Distribution')
plt.show()
```



```
[14]: import matplotlib.pyplot as plt

# Compute the length of each utterance
df_vis['utterance_length'] = df_vis['Utterance'].apply(len)

# Create a histogram
plt.figure(figsize=(8, 6))
plt.hist(df_vis['utterance_length'], bins=50)
plt.xlabel('Utterance Length')
plt.ylabel('Frequency')
plt.title('Distribution of Utterance Length')
plt.show()
```



```
[15]: def compute_lexical_diversity(text):
    # Tokenize the text into individual words
    tokens = nltk.word_tokenize(text)

    # Calculate the number of unique words (vocabulary)
    vocabulary_size = len(set(tokens))

    # Calculate the lexical diversity as the ratio of unique words to total
    ↪ words
    lexical_diversity = vocabulary_size / len(tokens) if len(tokens) > 0 else 0

    return lexical_diversity

# Compute the lexical diversity for each utterance
df_vis['lexical_diversity'] = df_vis['Utterance'].
    ↪ apply(compute_lexical_diversity)
```

```
[16]: def compute_word_frequency(text):
    # Tokenize the text into individual words
```

```

tokens = nltk.word_tokenize(text)

# Compute the word frequency distribution
word_freq = nltk.FreqDist(tokens)

return word_freq

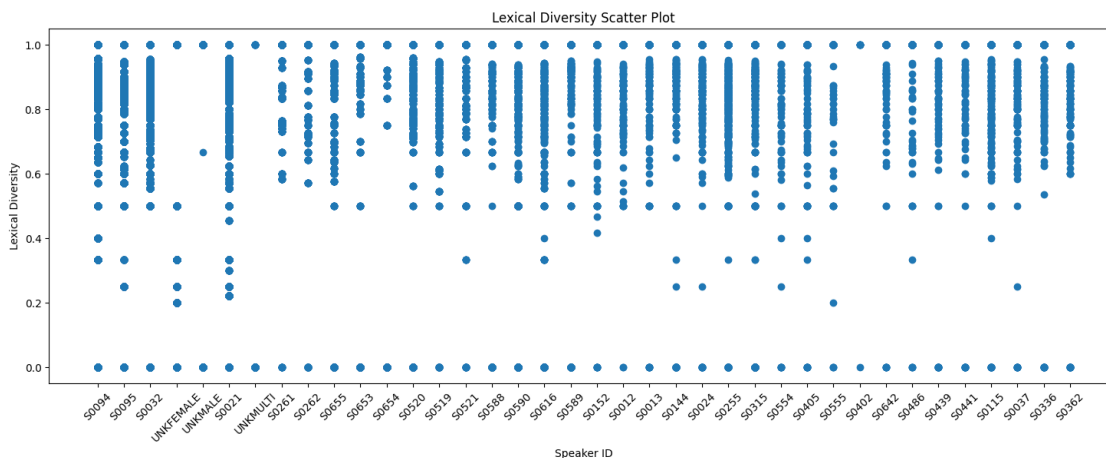
```

```

[17]: import matplotlib.pyplot as plt

# Create a scatter plot
plt.figure(figsize=(18, 6))
plt.scatter(df_vis['SpeakerID'], df_vis['lexical_diversity'])
plt.xlabel('Speaker ID')
plt.ylabel('Lexical Diversity')
plt.title('Lexical Diversity Scatter Plot')
plt.xticks(rotation=45)
plt.show()

```



```

[18]: import matplotlib.pyplot as plt
from nltk import FreqDist

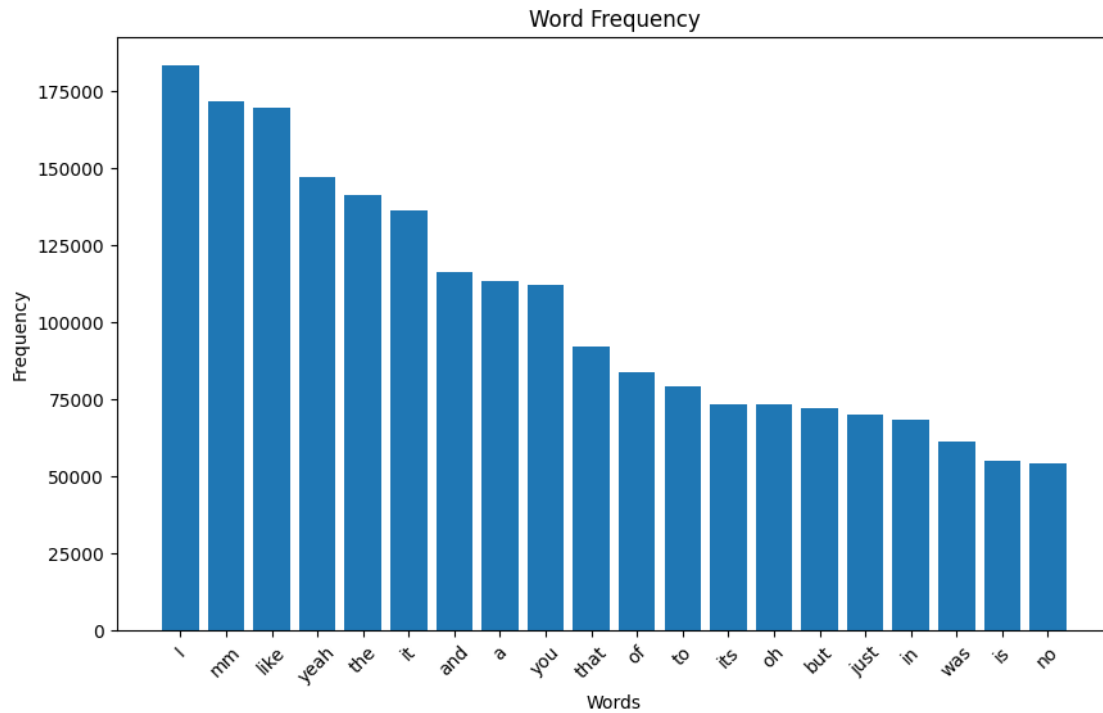
# Calculate word frequency
word_freq = compute_word_frequency(df_vis['Utterance'].str.cat(sep=' '))

# Get the most common words and their frequencies
most_common = word_freq.most_common(20)
words = [word[0] for word in most_common]
frequencies = [freq[1] for freq in most_common]

# Plot the word frequency
plt.figure(figsize=(10, 6))

```

```
plt.bar(words, frequencies)
plt.xlabel('Words')
plt.ylabel('Frequency')
plt.title('Word Frequency')
plt.xticks(rotation=45)
plt.show()
```



```
[19]: for index, row in df_vis.iterrows():
        utterance = row['Utterance']
        features = extract_linguistic_features(utterance)[1]
        for feature, value in features.items():
            df_vis.loc[index, feature] = value
```

```
[20]: list(extract_linguistic_features(utterance)[1].keys())
```


```
[20]: ['lexical_diversity',
        'sentence_length',
        'discourse_marker_count',
        'filler_count',
        'hesitation_count',
        'speech_rate']
```

```
[21]: import seaborn as sns
import matplotlib.pyplot as plt

# Select the columns containing the linguistic features
linguistic_feature_columns = list(extract_linguistic_features(utterance)[1].
    ↪keys())

# Calculate the correlation matrix
correlation_matrix = df_vis[linguistic_feature_columns].corr()

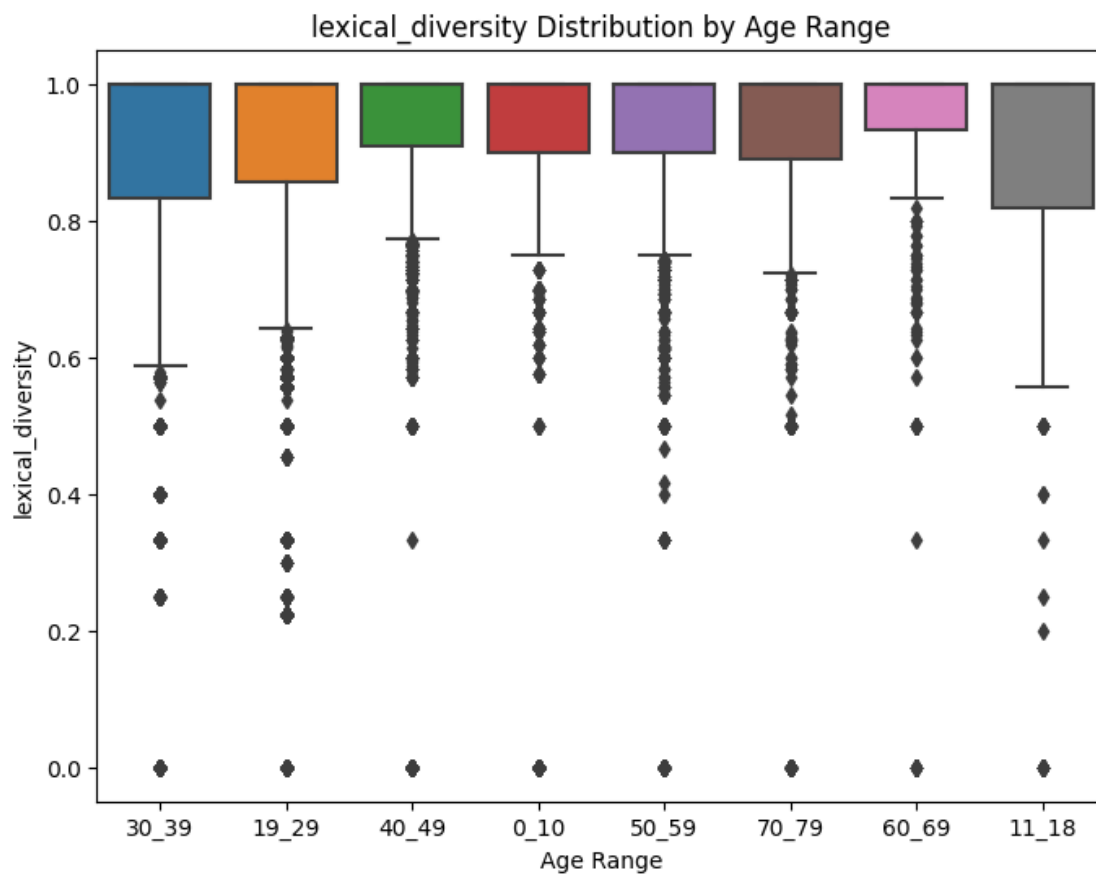
# Create the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', square=True)
plt.title('Linguistic Feature Correlation')
plt.show()
```



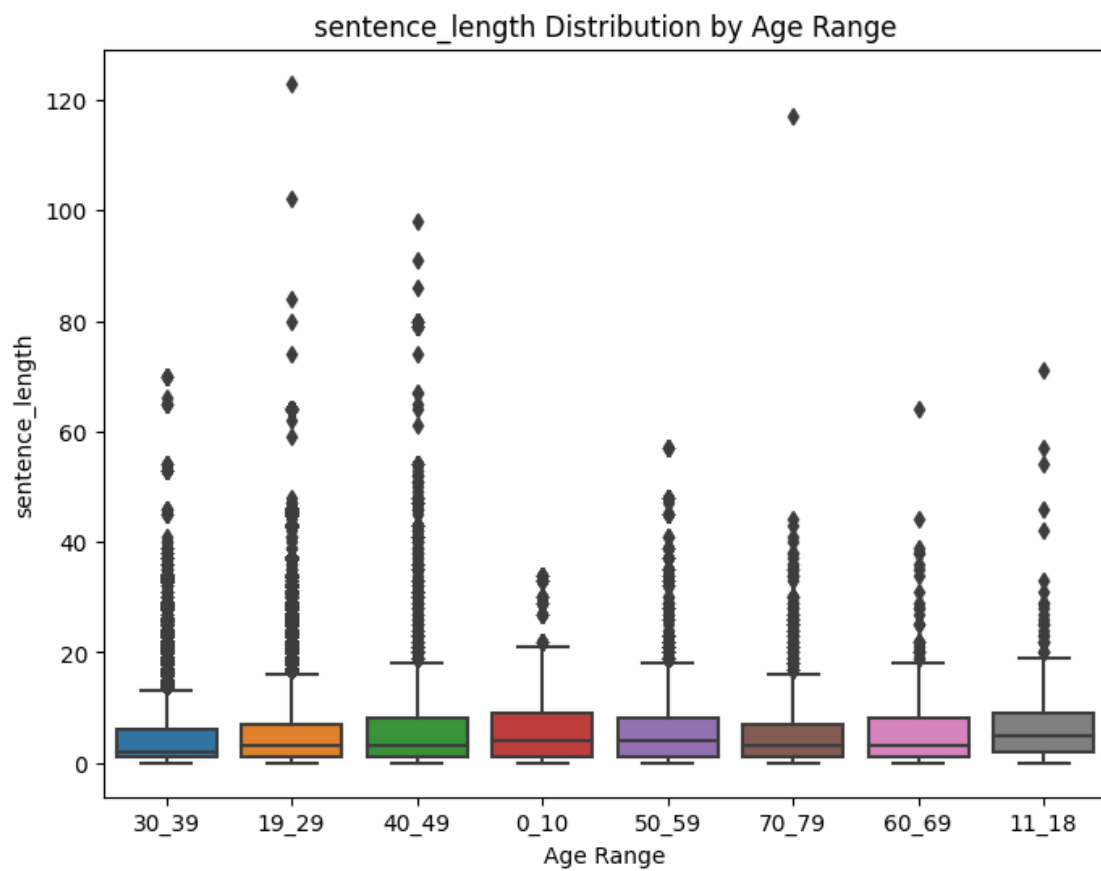
```
[22]: import seaborn as sns
import matplotlib.pyplot as plt

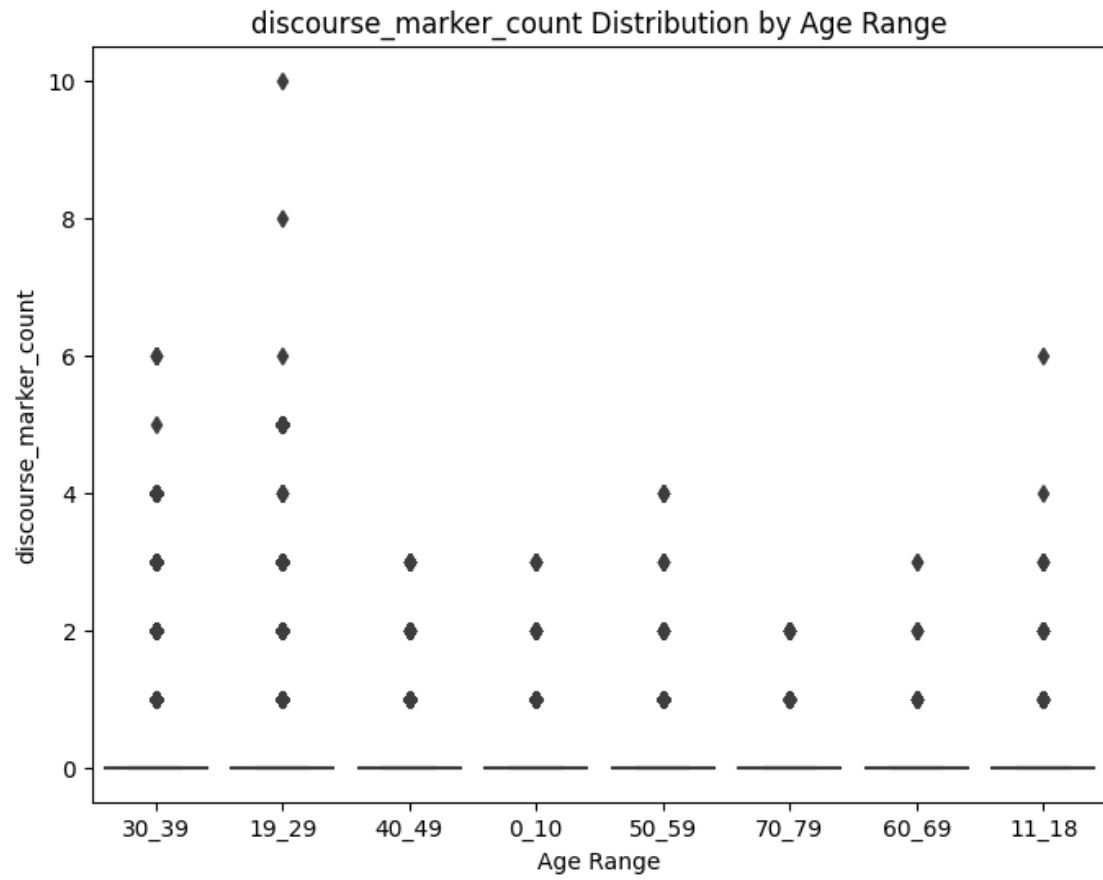
features = list(extract_linguistic_features(utterance)[1].keys())

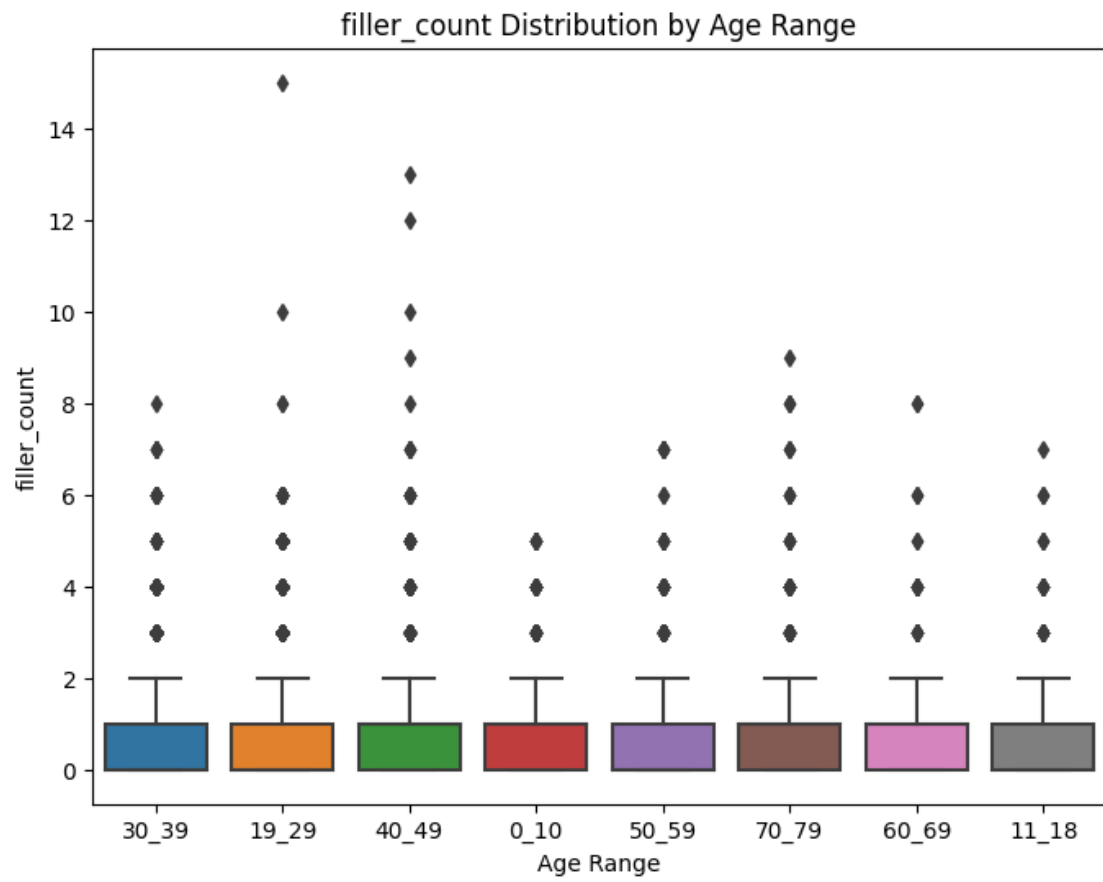
for feature in features:
    plt.figure(figsize=(8, 6))
    sns.boxplot(x='AgeLabel', y=feature, data=df_vis)
    plt.title(f'{feature} Distribution by Age Range')
    plt.xlabel('Age Range')
    plt.ylabel(feature)
    plt.show()
```

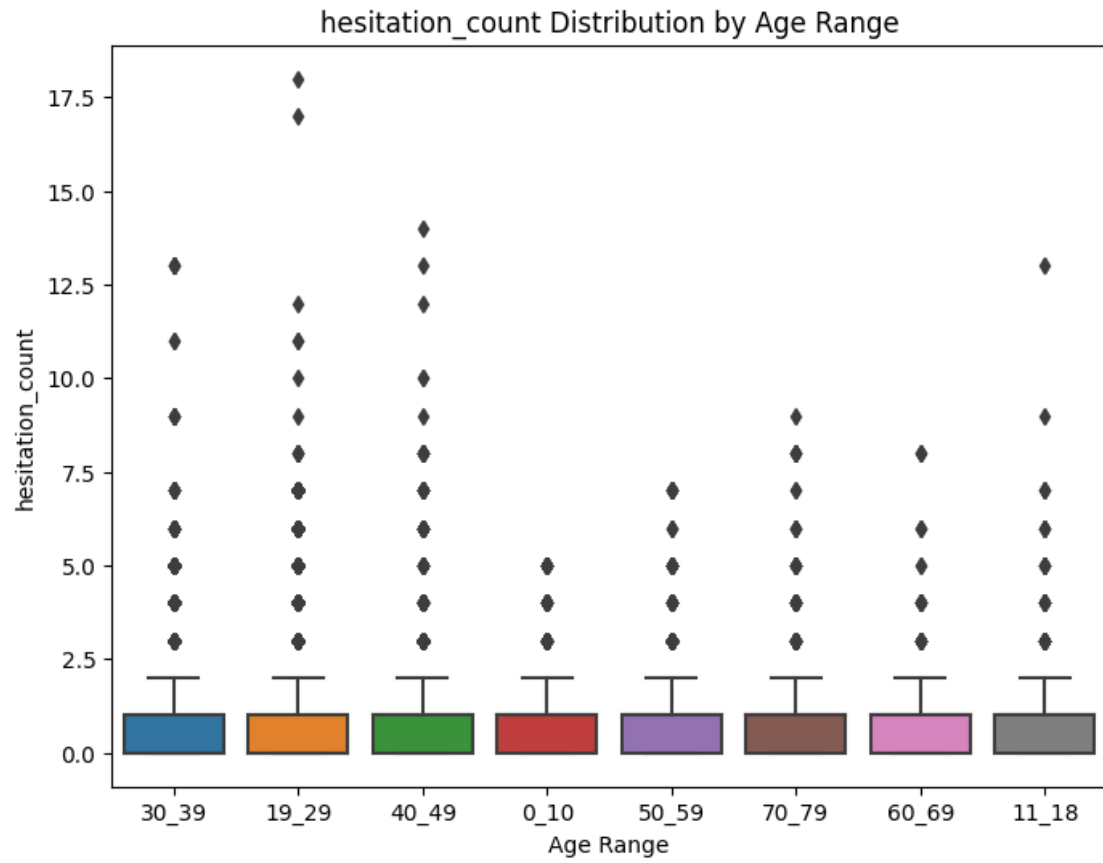


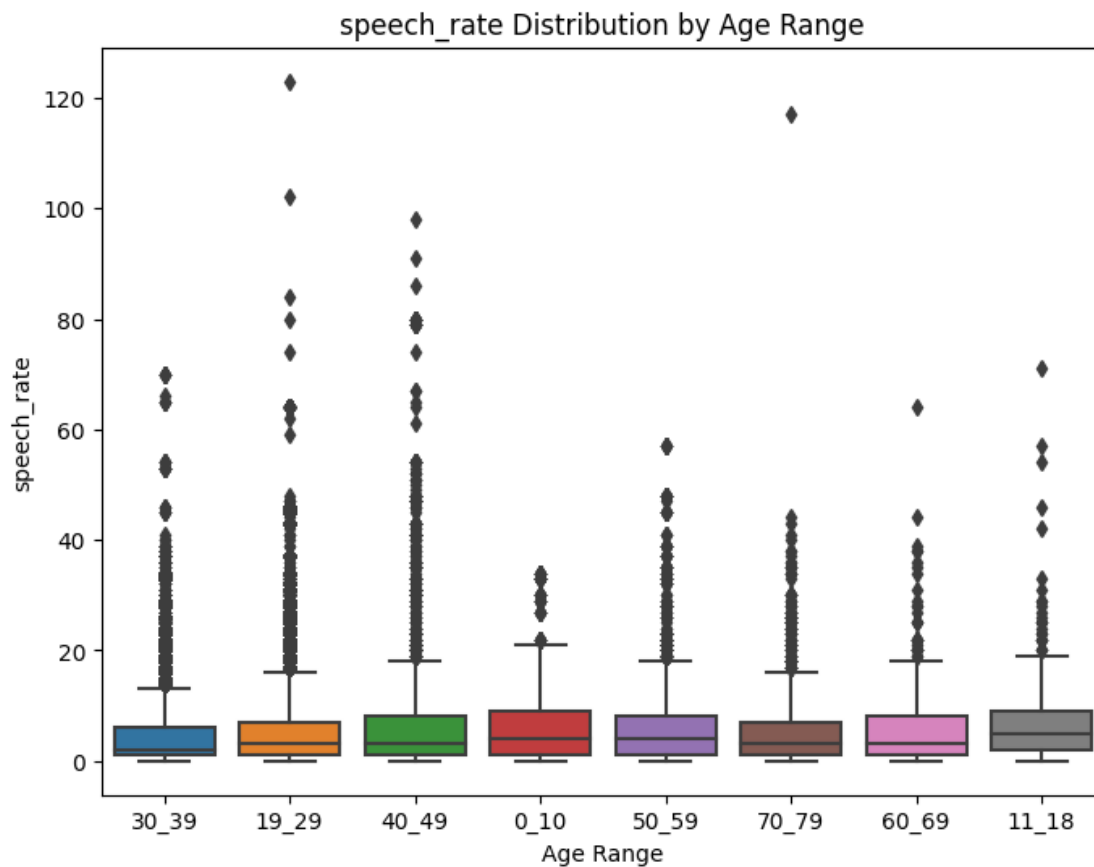












```
[23]: len(list(extract_linguistic_features(utterance)[1].keys()))
```

```
[23]: 6
```

```
[24]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Select the relevant features for clustering
selected_features = list(extract_linguistic_features(utterance)[1].keys())

# Extract the selected features from the dataset
X = df_vis[selected_features].values

# Scale or normalize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```

# Choose the number of clusters
num_clusters = 6

# Apply K-means clustering
kmeans = KMeans(n_clusters=num_clusters, random_state=42)
clusters = kmeans.fit_predict(X_scaled)

# Add the cluster labels to the dataframe
df_vis['Cluster'] = clusters


plt.figure(figsize=(20, 3))

# Visualize the clusters
plt.scatter(X[:, 0], X[:, 1], c=clusters)
for i, txt in enumerate(selected_features):
    plt.annotate(txt, (X[i, 0], X[i, 1]), fontsize=10)
plt.title('Clustering Results')
plt.show()

```

c:\Users\abhi\j\conda\envs\dmassessment\Lib\site-packages\sklearn\cluster\\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

```
warnings.warn(
```



```

[105]: # get the time variables from untagged and create a time-series analysis on
        ↳ dialogue data
df_time = df_all_utts.copy()

# drop unclear, trans, whoConfidence, vocal, foreign, anon, pause, trunc,
        ↳ shift, event
df_time.drop(['unclear', 'trans', 'whoConfidence', 'vocal', 'foreign', 'anon',
        ↳ 'pause', 'trunc', 'shift', 'event'], axis=1, inplace=True)

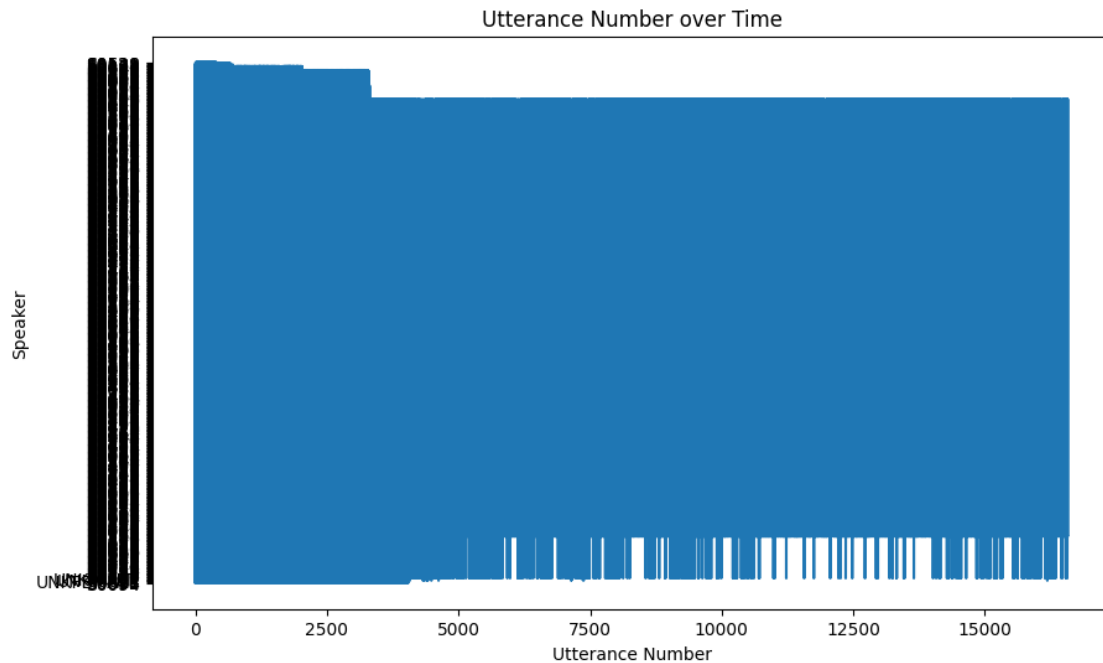
df_time.set_index('n', inplace=True)

print(df_time.columns)

```

```
# Plot the time-series data
plt.figure(figsize=(10, 6))
plt.plot(df_time['who'])
plt.title('Utterance Number over Time')
plt.xlabel('Utterance Number')
plt.ylabel("Speaker")
plt.show()
```

```
Index(['who', 'u'], dtype='object')
```



```
[25]: # get a copy of the dataframe
df_copy = df.copy()

# Replace age label values
df_copy['AgeLabel'].replace({'19_29': 'young', '50_59': 'old', '60_69': 'old', '70_79': 'old', '80_89': 'old', '90_99': 'old'}, inplace=True)

# Drop rows with remaining age labels
df_copy = df_copy[df_copy['AgeLabel'].isin(['young', 'old'])]

# remove all the None values from Utterance column
df_copy = df_copy.dropna(subset=['Utterance'])

df_copy
```

```
[25]:
```

	SpeakerID	AgeLabel	Utterance
2	S0032	young	yeah yeah \
4	S0032	young	I its something I
6	S0032	young	
10	S0032	young	
12	S0021	young	
...	...	...	...
1248104	S0032	young	no its not
1248105	S0021	young	I have like
1248107	S0032	young	
1248108	S0021	young	I have like tomato ketchup emergencies if I ru...
1248109	S0032	young	thats not as bad as the barbecue that we had t...

	Features
2	lexical_diversity:0.5 sentence_length:2.0 disc...
4	lexical_diversity:0.75 sentence_length:4.0 dis...
6	lexical_diversity:0 sentence_length:0 discours...
10	lexical_diversity:0 sentence_length:0 discours...
12	lexical_diversity:0 sentence_length:0 discours...
...	...
1248104	lexical_diversity:1.0 sentence_length:3.0 disc...
1248105	lexical_diversity:1.0 sentence_length:3.0 disc...
1248107	lexical_diversity:0 sentence_length:0 discours...
1248108	lexical_diversity:0.7391304347826086 sentence_...
1248109	lexical_diversity:0.9090909090909091 sentence_...

[607368 rows x 4 columns]

```
[26]: df_copy['AgeLabel'].unique()
```

```
[26]: array(['young', 'old'], dtype=object)
```

```
[27]: import pandas as pd
from sklearn.utils import resample

# Define the input features (X) and target variable (y)
X = df_copy[['Utterance']]
y = df_copy['AgeLabel']

# Combine X and y into a single DataFrame
df = pd.concat([X, y], axis=1)

# Separate the majority and minority classes
young_class = df[df['AgeLabel'] == 'young']
old_class = df[df['AgeLabel'] == 'old']

# Undersample the majority class
```



```

young_undersampled = resample(young_class, replace=False,
    ↪n_samples=len(old_class), random_state=42)

# Combine the undersampled majority class and the minority class
balanced_data = pd.concat([young_undersampled, old_class])

# Separate the features (X) and the target variable (y)
X_balanced = balanced_data[['Utterance']]
y_balanced = balanced_data['AgeLabel']

```

```
[28]: X_balanced.shape, y_balanced.shape
```

```
[28]: ((15528, 1), (15528,))
```

```
[29]: y_balanced.value_counts()
```

```

[29]: AgeLabel
      young      7764
      old       7764
      Name: count, dtype: int64

```

```

[30]: X_train, X_test, y_train, y_test = train_test_split(X_balanced, y_balanced,
    ↪test_size=0.2, random_state=42)

```

```
[31]: X_train.shape
```

```
[31]: (12422, 1)
```

```
[32]: X_test.shape
```

```
[32]: (3106, 1)
```

```
[33]: print(y_train.value_counts())
```

```

AgeLabel
old      6220
young    6202
      Name: count, dtype: int64

```

```

[34]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# Initialize the vectorizer (choose either CountVectorizer or TfidfVectorizer)
vectorizer = CountVectorizer() # or TfidfVectorizer()

# Fit the vectorizer on the training data
X_train_text = X_train['Utterance'] # assuming 'Utterance' column contains the
    ↪text data

```

```
X_train_vectorized = vectorizer.fit_transform(X_train_text)

# Transform the testing data using the fitted vectorizer
X_test_text = X_test['Utterance'] # assuming 'Utterance' column contains the
↳text data
X_test_vectorized = vectorizer.transform(X_test_text)
```

[35]: *# Step 6: Model Training with no Linguistic Features*

```
model = LogisticRegression()
model.fit(X_train_vectorized, y_train)
```

c:\Users\abhij\conda\envs\dmassessment\Lib\site-packages\sklearn\linear\_model\\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

[35]: LogisticRegression()

[36]: *# Step 7: Model Evaluation for old class*

```
y_pred = model.predict(X_test_vectorized)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, pos_label="old")
recall = recall_score(y_test, y_pred, pos_label="old")
f1 = f1_score(y_test, y_pred, pos_label="old")

print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
```

Accuracy: 0.8161622665808114

Precision: 0.8462633451957295

Recall: 0.7700777202072538

F1 Score: 0.8063750423872499

[37]: *# Step 7: Model Evaluation for young class*

```
y_pred = model.predict(X_test_vectorized)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, pos_label="young")
recall = recall_score(y_test, y_pred, pos_label="young")
```

```
f1 = f1_score(y_test, y_pred, pos_label="young")
```

```
print(f"Accuracy: {accuracy}")
```

```
print(f"Precision: {precision}")
```

```
print(f"Recall: {recall}")
```

```
print(f"F1 Score: {f1}")
```

Accuracy: 0.8161622665808114

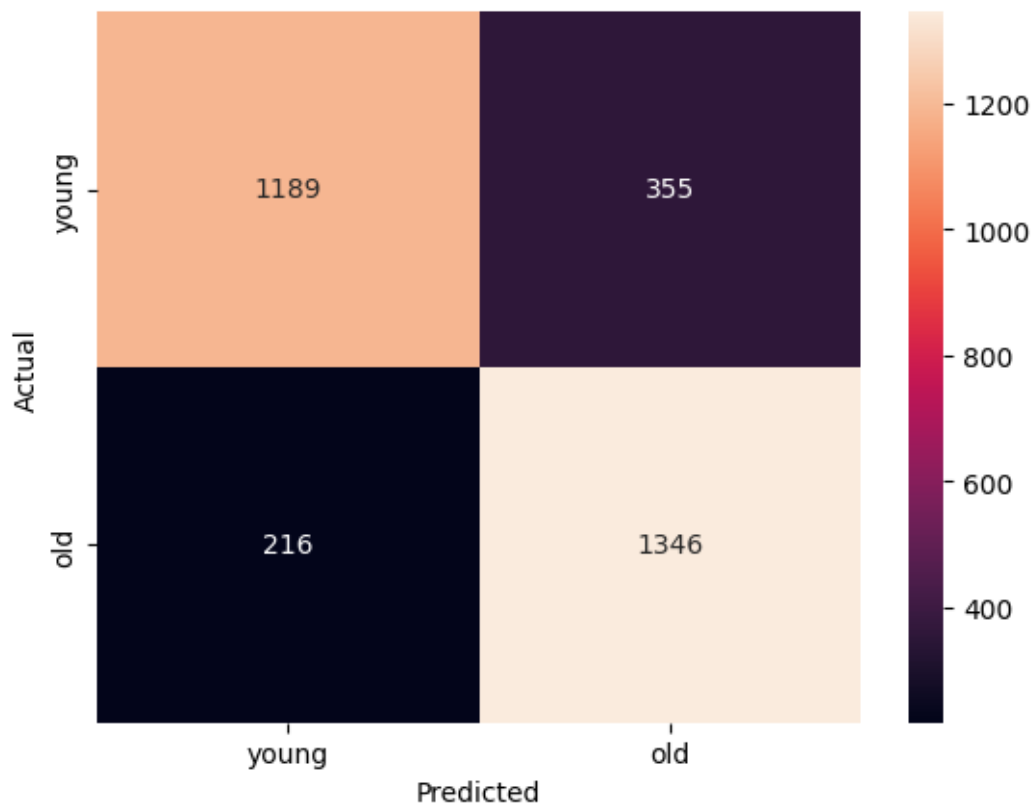
Precision: 0.7912992357436802

Recall: 0.8617157490396927

F1 Score: 0.8250076616610481

```
[38]: # plot confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns

conf_mat = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_mat, annot=True, fmt='d', xticklabels=['young', 'old'],
            yticklabels=['young', 'old'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```



```
[39]: df_vis.shape
```

```
[39]: (1248110, 12)
```

```
[40]: df_vis['AgeLabel'].value_counts()
```

```
[40]: AgeLabel
19_29    599604
30_39    547325
40_49     4896
50_59     4478
70_79     2818
0_10      2180
11_18      470
60_69      468
Name: count, dtype: int64
```

```
[41]: # get a copy of the dataframe
df_with_features = df_vis.copy()

# convert the AgeLabel column, 19_29 to young and 50_59, 60_69, 70_79, 80_89,
↳ 90_99 to old
df_with_features['AgeLabel'].replace({'19_29': 'young', '50_59': 'old', '60_69':
↳ 'old', '70_79': 'old', '80_89': 'old', '90_99': 'old'}, inplace=True)

# drop rows with remaining age labels
df_with_features = df_with_features[df_with_features['AgeLabel'].isin(['young',
↳ 'old'])]
```

```
[42]: df_with_features['AgeLabel'].value_counts()
```

```
[42]: AgeLabel
young    599604
old       7764
Name: count, dtype: int64
```

```
[43]: df_with_features
```

```
[43]:
```

	SpeakerID	AgeLabel	Utterance
2	S0032	young	yeah yeah \
4	S0032	young	I its something I
6	S0032	young	
10	S0032	young	
12	S0021	young	
...	...	...	...
1248104	S0032	young	no its not

1248105	S0021	young	I have like
1248107	S0032	young	
1248108	S0021	young	I have like tomato ketchup emergencies if I ru...
1248109	S0032	young	thats not as bad as the barbecue that we had t...

	Features	utterance_length
2	lexical_diversity:0.5 sentence_length:2.0 disc...	9 \
4	lexical_diversity:0.75 sentence_length:4.0 dis...	17
6	lexical_diversity:0 sentence_length:0 discours...	0
10	lexical_diversity:0 sentence_length:0 discours...	0
12	lexical_diversity:0 sentence_length:0 discours...	0
...	...	...
1248104	lexical_diversity:1.0 sentence_length:3.0 disc...	10
1248105	lexical_diversity:1.0 sentence_length:3.0 disc...	11
1248107	lexical_diversity:0 sentence_length:0 discours...	0
1248108	lexical_diversity:0.7391304347826086 sentence_...	119
1248109	lexical_diversity:0.9090909090909091 sentence_...	51

	lexical_diversity	sentence_length	discourse_marker_count
2	0.500000	2.0	0.0 \
4	0.750000	4.0	0.0
6	0.000000	0.0	0.0
10	0.000000	0.0	0.0
12	0.000000	0.0	0.0
...	...	...	...
1248104	1.000000	3.0	0.0
1248105	1.000000	3.0	1.0
1248107	0.000000	0.0	0.0
1248108	0.739130	23.0	2.0
1248109	0.909091	11.0	0.0

	filler_count	hesitation_count	speech_rate	Cluster
2	2.0	2.0	2.0	3
4	0.0	0.0	4.0	1
6	0.0	0.0	0.0	0
10	0.0	0.0	0.0	0
12	0.0	0.0	0.0	0
...	...	...	...	...
1248104	0.0	0.0	3.0	1
1248105	0.0	1.0	3.0	4
1248107	0.0	0.0	0.0	0
1248108	2.0	4.0	23.0	2
1248109	0.0	0.0	11.0	1

[607368 rows x 12 columns]

```
[44]: df_with_features.columns
```

```
[44]: Index(['SpeakerID', 'AgeLabel', 'Utterance', 'Features', 'utterance_length',
        'lexical_diversity', 'sentence_length', 'discourse_marker_count',
        'filler_count', 'hesitation_count', 'speech_rate', 'Cluster'],
        dtype='object')
```

```
[45]: df_with_features.describe()
```

```
[45]:
```

	utterance_length	lexical_diversity	sentence_length	
count	607368.000000	607368.000000	607368.000000	\
mean	24.330457	0.810971	5.234224	
std	31.979818	0.357185	6.622813	
min	0.000000	0.000000	0.000000	
25%	2.000000	0.857143	1.000000	
50%	13.000000	1.000000	3.000000	
75%	35.000000	1.000000	7.000000	
max	630.000000	1.000000	123.000000	

	discourse_marker_count	filler_count	hesitation_count	speech_rate	
count	607368.000000	607368.000000	607368.000000	607368.000000	\
mean	0.169834	0.519395	0.689230	5.234224	
std	0.498674	0.737525	0.975447	6.622813	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	1.000000	
50%	0.000000	0.000000	0.000000	3.000000	
75%	0.000000	1.000000	1.000000	7.000000	
max	10.000000	15.000000	18.000000	123.000000	

	Cluster
count	607368.000000
mean	2.015722
std	1.481856
min	0.000000
25%	1.000000
50%	1.000000
75%	3.000000
max	5.000000

```
[46]: # get X and y
X = df_with_features[['Utterance', 'utterance_length', 'lexical_diversity',
    ↪ 'sentence_length', 'discourse_marker_count', 'filler_count',
    ↪ 'hesitation_count', 'speech_rate']]
y = df_with_features['AgeLabel']

# balance the dataset
young_class = df_with_features[df_with_features['AgeLabel'] == 'young']
old_class = df_with_features[df_with_features['AgeLabel'] == 'old']
```

```

young_undersampled = resample(young_class, replace=False,
    ↪n_samples=len(old_class), random_state=42)

balanced_data = pd.concat([young_undersampled, old_class])

# separate the features (X) and the target variable (y)
X_balanced = balanced_data[['Utterance', 'utterance_length',
    ↪'lexical_diversity', 'sentence_length', 'discourse_marker_count',
    ↪'filler_count', 'hesitation_count', 'speech_rate']]
y_balanced = balanced_data['AgeLabel']

X_balanced.shape, y_balanced.shape

```

```
[46]: ((15528, 8), (15528,))
```

```

[47]: # split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_balanced, y_balanced,
    ↪test_size=0.2, random_state=42)

```

```
[48]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[48]: ((12422, 8), (3106, 8), (12422,), (3106,))
```

```
[49]: y_train.value_counts()
```

```

[49]: AgeLabel
old      6220
young    6202
Name: count, dtype: int64

```

```

[50]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# Initialize the vectorizer (choose either CountVectorizer or TfidfVectorizer)
vectorizer = CountVectorizer() # or TfidfVectorizer()

# Fit the vectorizer on the training data
X_train_text = X_train['Utterance'] # assuming 'Utterance' column contains the
    ↪text data
X_train_vectorized = vectorizer.fit_transform(X_train_text)

# Transform the testing data using the fitted vectorizer
X_test_text = X_test['Utterance'] # assuming 'Utterance' column contains the
    ↪text data
X_test_vectorized = vectorizer.transform(X_test_text)

```

```
[51]: # Step 6: Model Training with Linguistic Features
model2 = LogisticRegression(max_iter=1000)
model2.fit(X_train_vectorized, y_train)
```

```
[51]: LogisticRegression(max_iter=1000)
```

```
[52]: # Step 7: Model Evaluation
from sklearn.metrics import classification_report

y_pred = model2.predict(X_test_vectorized)


print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
old	0.85	0.77	0.81	1544
young	0.79	0.86	0.83	1562
accuracy			0.82	3106
macro avg	0.82	0.82	0.82	3106
weighted avg	0.82	0.82	0.82	3106

```
[53]: # plot confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns

conf_mat = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_mat, annot=True, fmt='d', xticklabels=['young', 'old'],
            yticklabels=['young', 'old'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```





By comparing both models, one with just utterances and the utterances and features, we expect to see a difference in the performance of the models. But in this case, the performance of the model is almost the same. I assume this is because the data is not large enough to make a difference in the performance of the model. If we had a large dataset, then we would have seen a difference in the performance of the model.

```
[54]: # predictions for the first 5 test samples

X_test_transformed = vectorizer.inverse_transform(X_test_vectorized)

predictions = model.predict(X_test_vectorized)

# for i in range(len(X_test_transformed)):
for i in range(5):
    sentence = " ".join(X_test_transformed[i])
    predicted_output = predictions[i]
    real_output = y_test.iloc[i]

    print("Input Sentence: ", sentence)
    print("Predicted Output: ", predicted_output)
    print("Real Output: ", real_output)
```

```
print()
```

```
Input Sentence:  he
Predicted Output: old
Real Output:  young
```

```
Input Sentence:  and need no pool swimming the thing whole you
Predicted Output: old
Real Output:  old
```

```
Input Sentence:
Predicted Output:  young
Real Output:  young
```

```
Input Sentence:  mm
Predicted Output:  young
Real Output:  young
```

```
Input Sentence:  mm
Predicted Output:  young
Real Output:  old
```

```
[55]: # get the predictions for the whole test set and display the results in a
      ↪dataframe
X_test_transformed = vectorizer.inverse_transform(X_test_vectorized)
predictions = model.predict(X_test_vectorized)

results_df = pd.DataFrame({'Input Sentence': [' ' .join(sentence) for sentence_
      ↪in X_test_transformed],
                           'Predicted Output': predictions,
                           'Real Output': y_test.values})

results_df
```

```
[55]:
```

	Input Sentence	Predicted Output
0	he	old \
1	and need no pool swimming the thing whole you	old
2		young
3	mm	young
4	mm	young
...	...	...
3101	yeah	old
3102	is it so time whatever	young
3103	did make these you	young
3104	ill speak to	old
3105	across did er just oh the to	old

	Real Output
0	young
1	old
2	young
3	young
4	old
...	...
3101	old
3102	young
3103	old
3104	old
3105	old

[3106 rows x 3 columns]

```
[56]: wrong_predictions = results_df[results_df['Predicted Output'] !=
    ↪results_df['Real Output']]
wrong_predictions
```

```
[56]:
```

	Input Sentence	Predicted Output	Real Output
0	he	old	young
4	mm	young	old
11	yeah	old	young
18	an day eight hour thats	young	old
19	er	old	young
...	...	...	...
3083	didnt help it last	young	old
3088	so	young	old
3096	is it yeah	old	young
3098	where	young	old
3103	did make these you	young	old

[571 rows x 3 columns]

```
[57]: # # Get feature names
# import numpy as np

# # Retrieve the coefficients from the trained model
# coefficients = model2.coef_[0]

# # Get feature names from the vectorizer
# feature_names = vectorizer.get_feature_names()

# print(len(feature_names))

# # Get feature names
```

```

# feature_names = ['Utterance', 'utterance_length', 'lexical_diversity',
↳ 'sentence_length',
#                 'discourse_marker_count', 'filler_count', 'hesitation_count',
#                 'speech_rate']

# # Plot the feature importances
# plt.figure(figsize=(10, 6))
# plt.barh(range(len(feature_names)), np.abs(coefficients), align='center')
# plt.yticks(range(len(feature_names)), feature_names)
# plt.xlabel('Coefficient Magnitude')
# plt.ylabel('Features')
# plt.title('Feature Importance for Logistic Regression')
# plt.tight_layout()
# plt.show()

```

[58]: # predictions for the first 5 test samples

```

X_test_transformed = vectorizer.inverse_transform(X_test_vectorized)

predictions = model.predict(X_test_vectorized)

# for i in range(len(X_test_transformed)):
for i in range(5):
    sentence = " ".join(X_test_transformed[i])
    predicted_output = predictions[i]
    real_output = y_test.iloc[i]

    print("Input Sentence: ", sentence)
    print("Predicted Output: ", predicted_output)
    print("Real Output: ", real_output)
    print()

```

Input Sentence: he  
Predicted Output: old  
Real Output: young

Input Sentence: and need no pool swimming the thing whole you  
Predicted Output: old  
Real Output: old

Input Sentence:  
Predicted Output: young  
Real Output: young

Input Sentence: mm  
Predicted Output: young  
Real Output: young

Input Sentence: mm  
 Predicted Output: young  
 Real Output: old

```
[59]: # get the predictions for the whole test set and display the results in a
      ↪ dataframe
X_test_transformed = vectorizer.inverse_transform(X_test_vectorized)
predictions = model.predict(X_test_vectorized)

results_df = pd.DataFrame({'Input Sentence': [' '.join(sentence) for sentence
      ↪ in X_test_transformed],
                           'Predicted Output': predictions,
                           'Real Output': y_test.values})

results_df
```

```
[59]:
```

	Input Sentence	Predicted Output
0	he	old \
1	and need no pool swimming the thing whole you	old
2		young
3	mm	young
4	mm	young
...	...	...
3101	yeah	old
3102	is it so time whatever	young
3103	did make these you	young
3104	ill speak to	old
3105	across did er just oh the to	old

	Real Output
0	young
1	old
2	young
3	young
4	old
...	...
3101	old
3102	young
3103	old
3104	old
3105	old

[3106 rows x 3 columns]

```
[60]: wrong_predictions = results_df[results_df['Predicted Output'] !=
      ↪results_df['Real Output']]
      wrong_predictions
```

```
[60]:
```

	Input Sentence	Predicted Output	Real Output
0	he	old	young
4	mm	young	old
11	yeah	old	young
18	an day eight hour thats	young	old
19	er	old	young
...	...	...	...
3083	didnt help it last	young	old
3088	so	young	old
3096	is it yeah	old	young
3098	where	young	old
3103	did make these you	young	old

```
[571 rows x 3 columns]
```

## 2 Using Bert Model

```
[61]: df.columns
```

```
[61]: Index(['Utterance', 'AgeLabel'], dtype='object')
```

```
[62]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from transformers import BertTokenizer, BertForSequenceClassification, AdamW
      import torch

      # Split the data into training and testing sets
      train_df, test_df = train_test_split(df, test_size=0.2, random_state=42)

      # Get the training and testing utterances and labels
      train_texts = train_df['Utterance'].tolist()
      train_labels = train_df['AgeLabel'].tolist()
      test_texts = test_df['Utterance'].tolist()
      test_labels = test_df['AgeLabel'].tolist()
```

```
c:\Users\abhiij\.conda\envs\dmassessment\Lib\site-packages\tqdm\auto.py:21:
TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
      from .autonotebook import tqdm as notebook_tqdm
```

```
[63]: import numpy as np
```

```
train_labels = np.array(train_labels)
test_labels = np.array(test_labels)
```

```
[64]: # Load the BERT tokenizer
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

# Load the BERT model for sequence classification
model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
↳ num_labels=2)
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.seq\_relationship.weight', 'cls.predictions.transform.LayerNorm.weight', 'cls.seq\_relationship.bias', 'cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.transform.dense.bias']

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
[66]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
train_labels = label_encoder.fit_transform(train_labels)
test_labels = label_encoder.transform(test_labels)
```

```
[89]: test_labels
```

```
[89]: array([1, 1, 1, ..., 1, 1, 1])
```

```
[67]: # Tokenize and encode the training utterances
train_encodings = tokenizer(train_texts, truncation=True, padding=True)

# Tokenize and encode the testing utterances
test_encodings = tokenizer(test_texts, truncation=True, padding=True)

# Create torch tensors for the encoded data and labels
```

```

train_dataset = torch.utils.data.TensorDataset(
    torch.tensor(train_encodings['input_ids']),
    torch.tensor(train_encodings['attention_mask']),
    torch.tensor(train_labels)
)
test_dataset = torch.utils.data.TensorDataset(
    torch.tensor(test_encodings['input_ids']),
    torch.tensor(test_encodings['attention_mask']),
    torch.tensor(test_labels)
)

```

```

[68]: # Define the training parameters
batch_size = 16
epochs = 3

# Create a DataLoader for the training and testing datasets
train_loader = torch.utils.data.DataLoader(train_dataset,
    ↪ batch_size=batch_size, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size)

```

```

[69]: # import torch
# import torch.cuda

# # Set the max_split_size_mb parameter
# torch.cuda.set_per_process_memory_fraction(fraction=0.2)

# torch.cuda.empty_cache()

```

```

[70]: # import gc

# gc.collect()

# torch.cuda.empty_cache()

```

```

[71]: # # Set the maximum memory split size in megabytes
# torch.backends.cuda.max_split_size_mb = 2000 # Adjust the value as needed

```

```

[72]: # import torch
# from GPUtil import showUtilization as gpu_usage
# from numba import cuda

# def free_gpu_cache():
#     print("Initial GPU Usage")
#     gpu_usage()

#     torch.cuda.empty_cache()

```



```

#     cuda.select_device(0)
#     cuda.close()
#     cuda.select_device(0)

#     print("GPU Usage after emptying the cache")
#     gpu_usage()

# free_gpu_cache()

```

```
[73]: # gpu_usage()
```

```
[74]: # torch.cuda.device_count()
# cuda.select_device(0)
```

```
[75]: # import os
# os.environ['CUDA_LAUNCH_BLOCKING'] = '1'
```

```
[76]: # torch.cuda.memory_allocated()
```

```
[77]: # Set the device (CPU or GPU)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Move the model to the device
model.to(device)

# Set the optimizer and learning rate
optimizer = AdamW(model.parameters(), lr=1e-5)

# Training loop
model.train()
for epoch in range(epochs):
    for batch in train_loader:
        input_ids, attention_mask, labels = batch
        print("Input ids:", input_ids.shape, "Attention mask:", attention_mask.
↪shape, "Labels:", labels.shape)
        input_ids = input_ids.to(device)
        attention_mask = attention_mask.to(device)
        labels = labels.to(device)
        labels = labels.to(torch.long)

        optimizer.zero_grad()

        outputs = model(input_ids, attention_mask=attention_mask, labels=labels)
        loss = outputs.loss
        loss.backward()
        optimizer.step()

```

```

torch.Size([16])
Input ids: torch.Size([16, 139]) Attention mask: torch.Size([16, 139]) Labels:
torch.Size([16])
Input ids: torch.Size([16, 139]) Attention mask: torch.Size([16, 139]) Labels:
torch.Size([16])
Input ids: torch.Size([16, 139]) Attention mask: torch.Size([16, 139]) Labels:
torch.Size([16])
Input ids: torch.Size([16, 139]) Attention mask: torch.Size([16, 139]) Labels:
torch.Size([16])
Input ids: torch.Size([16, 139]) Attention mask: torch.Size([16, 139]) Labels:
torch.Size([16])
Input ids: torch.Size([6, 139]) Attention mask: torch.Size([6, 139]) Labels:
torch.Size([6])

```

```

[83]: # save the model
      model.save_pretrained('models/bert_model')

```

```

[84]: # get model from saved
      model = BertForSequenceClassification.from_pretrained('models/bert_model',
      ↪num_labels=2)

      device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

      model.to(device)

      # evaluate the model
      model.eval()
      with torch.no_grad():
          correct = 0
          total = 0
          for batch in test_loader:
              input_ids, attention_mask, labels = batch
              input_ids = input_ids.to(device)
              attention_mask = attention_mask.to(device)
              labels = labels.to(device)

              outputs = model(input_ids, attention_mask=attention_mask)
              _, predicted = torch.max(outputs.logits, dim=1)

              total += labels.size(0)
              correct += (predicted == labels).sum().item()

          accuracy = correct / total
          print('Accuracy on test data: {:.2f}%'.format(accuracy * 100))

```

Accuracy on test data: 99.47%

```
[85]: # Confusion matrix
from sklearn.metrics import confusion_matrix

model.eval()
y_pred = []
y_true = []
with torch.no_grad():
    for batch in test_loader:
        input_ids, attention_mask, labels = batch
        input_ids = input_ids.to(device)
        attention_mask = attention_mask.to(device)
        labels = labels.to(device)

        outputs = model(input_ids, attention_mask=attention_mask)
        _, predicted = torch.max(outputs.logits, dim=1)

        y_pred.extend(predicted.tolist())
        y_true.extend(labels.tolist())

cm = confusion_matrix(y_true, y_pred)
print(cm)
```

```
[[ 997   572]
 [  71 119834]]
```

```
[87]: # get all the predictions for test data
model.eval()
y_pred = []
y_true = []
X_input = []
with torch.no_grad():
    for batch in test_loader:
        input_ids, attention_mask, labels = batch
        input_ids = input_ids.to(device)
        attention_mask = attention_mask.to(device)
        labels = labels.to(device)

        outputs = model(input_ids, attention_mask=attention_mask)
        _, predicted = torch.max(outputs.logits, dim=1)

        X_input.extend(input_ids.tolist())
        y_pred.extend(predicted.tolist())
        y_true.extend(labels.tolist())

# create a dataframe with the predictions
bert_results_df = pd.DataFrame({'text': test_texts, 'label': y_true,
↪ 'prediction': y_pred})
```

```
[91]: # convert the label and prediction to the actual class names
bert_results_df['label'] = bert_results_df['label'].apply(lambda x: 'young' if
↳ x == 1 else 'old')
bert_results_df['prediction'] = bert_results_df['prediction'].apply(lambda x:
↳ 'young' if x == 1 else 'old')

bert_results_df
```

```
[91]:
```

	text	label	prediction
0	that wed be able to drink together no	young	young
1	yeah yeah coming	young	young
2	think we drank it all though didnt we	young	young
3		young	young
4	yeah all good thank you	young	young
...	...	...	...
121469	mm	young	young
121470	mm	young	young
121471	erm	young	young
121472	for Christmas	young	young
121473	suppose theyre made out of paper	young	young

[121474 rows x 3 columns]

```
[92]: # get all the wrong predictions
wrong_predictions_df = bert_results_df[bert_results_df['label'] !=
↳ bert_results_df['prediction']]
wrong_predictions_df
```

```
[92]:
```

	text	label	prediction
134	right	old	young
458	okay	old	young
762	yeah	old	young
1517	mm	old	young
1579	yeah	old	young
...	...	...	...
120317	yeah	old	young
120515	oh right yeah	old	young
120667	yeah	old	young
120853		old	young
121406		old	young

[643 rows x 3 columns]